

# COSE474 Deep Learning: Project #1: MLP Implementation

## Project report REUTELSTERZ ELIAS

### Description of the code:

In the first code fragment, I initialized and computed the different components of the 2-layer neural network ( $z_1$ ,  $p_1$ ,  $z_2$ ). I used the same notation as in our assignment and script. Because we use  $N$  different data points in each batch-iteration I decided to generate  $z_1$ ,  $p_1$ ,  $z_2$  as  $N$  rows with one row for every datapoint.

In the second section, the data loss and the two L2 regularizations were calculated. Here one could discuss, that in  $p_2$  we only need the entries  $p_2[i, y[i]]$  but later in the calculation of the gradient of  $p_2$  we also need the whole matrix. The total loss function contains the data loss and the regularization loss parts.

In the backpropagation section, I calculated the gradients for all data points and finally averaged them. To do this, I first had to calculate the four different gradients ( $b_2$  and  $W_2$  we already knew from the 2nd assignment). I calculated them for all data points in the order  $b_2 \rightarrow W_2 \rightarrow b_1 \rightarrow W_1$ , so that I can use the intermediate results for the next calculation. Finally, I had to add the derivative of the regularization terms to the gradients of  $W_1$  and  $W_2$  and store everything in the map `grads["xxx"]`.

In the function `train`, first a randomized selection of the training data points with the size `batch_size` had to be determined. Then the parameters were updated.

Finally, in the `predict` function, we were able to call the loss function without argument  $y$  to calculate the most likely labels, since it uses our final parameters to determine the scores for the labels.

### Results:

Because the validation accuracy was just as high as the training accuracy, I tried to test out how many hidden layers and iterations I can use without letting the gap be too big and still gaining test accuracy. I also increased the learning rate, so that we get away of the linear fitting and get an exponential fitting. I Additionally decreased the `batch_size` (but increased the iteration count in a higher relative relation) and it worked out for a better performance.

With trial and error in the fine-tuning I managed to get the rates:

- Training accuracy: 0.622
- Validation accuracy: 0.464
- Test accuracy: 0.476

while using the hyperparameters:

- `hidden_size` = 80
- `num_iters` = 5000
- `batch_size` = 50
- `learning_rate` =  $1e-3$
- `learning_rate_decay` = 0.95
- `reg` = 0.2

### Discussion:

While working on the project, I encountered two problems whose solution strategies I had to weigh. On the one hand, I tried to solve as many calculations as possible with direct matrix calculation to get the best performance. This was difficult, because I first calculated the gradients with the individual functions and could not always insert the case distinctions into the Numpy formula. Also, in some cases I find it more descriptive to go through the rows and columns one at a time and calculate the functions one at a time. In the end I managed to get rid of all for-loops to get a good performance.

The second problem was tuning the model. Here it is difficult to develop a feeling for which change, which consequences. For example, how high the deviation between training accuracy and test accuracy should be optimally, without getting an over- or underfitted model.